

# Report HW#3

## Team Member

Pawel Urbanowicz	108015016
Martin Ledl	108012012
Tobias Kick	108998413

## Github Repository

The Github repository for this project can be found here:

[https://github.com/mledl/BDMA\\_HW/tree/master/HW3](https://github.com/mledl/BDMA_HW/tree/master/HW3)

Please check the repository for results under “/data/results/”, because the limited upload size restricted us from uploading the result files to cyberclassroom.

## Responsibilities

### Pawel Urbanowicz

- Configure the Production Spark Environment and run the final solution for benchmark.
- Take care of task 2 (Min-Hashing)
- Contribute to report.

### Martin Ledl

- Setup project to work on HW#3.
- Take care of task 3(Locality-Sensitive Hashing)
- Contribute to report.

### Tobias Kick

- Take care of task 1(shingling)
- Optimizing different aspects of data processing in all 3 tasks
- Contribute to report.

## Environment Setup

For local development we tested our code on a locally installed spark instance and for target stage we used Docker technology to wrap spark master instance, 2 spark workers instances and our Python script into separate containers. Software/Frameworks in use:

- Python 3.7 to write our code
- Spark version 2.4.4
- Hadoop version 3.2.1

- Docker engine version 18.09.2.

## Environment setup for OSX

### a. Install Spark

brew install apache-spark

```
MBP-Pawel:~ pawelurbanowicz$ brew info apache-spark
apache-spark: stable 2.4.4, HEAD
Engine for large-scale data processing
https://spark.apache.org/
/usr/local/Cellar/apache-spark/2.4.0 (1,215 files, 249MB) *
  Built from source on 2019-03-20 at 02:46:21
From: https://github.com/Homebrew/homebrew-core/blob/master/Formula/apache-spark.rb
=> Requirements
Required: java = 1.8 ✓
=> Options
--HEAD
    Install HEAD version
=> Analytics
install: 5,390 (30 days), 15,259 (90 days), 62,289 (365 days)
install_on_request: 5,237 (30 days), 14,816 (90 days), 59,600 (365 days)
build_error: 0 (30 days)
```

### b. Install Hadoop

```
MBP-Pawel:~ pawelurbanowicz$ brew info hadoop
hadoop: stable 3.2.1
Framework for distributed processing of large data sets
https://hadoop.apache.org/
Conflicts with:
  yarn (because both install `yarn` binaries)
/usr/local/Cellar/hadoop/3.2.1 (22,397 files, 815.6MB)
  Built from source on 2019-10-15 at 17:58:46
From: https://github.com/Homebrew/homebrew-core/blob/master/Formula/hadoop.rb
=> Requirements
Required: java >= 1.8 ✓
=> Analytics
install: 4,381 (30 days), 10,643 (90 days), 44,685 (365 days)
install_on_request: 3,670 (30 days), 9,017 (90 days), 38,145 (365 days)
build_error: 0 (30 days)
```

### c. Install Docker Desktop for Mac

<https://docs.docker.com/docker-for-mac/install/>

```
MBP-Pawel:~ pawelurbanowicz$ docker --version
Docker version 18.09.2, build 6247962
MBP-Pawel:~ pawelurbanowicz$
```

## d. Install Docker Compose

```
brew install docker-compose
```

```

docker version 18.09.2, build 6247962
MBP-Pawel:~ pawelurbanowicz$ brew info docker-compose
docker-compose: stable 1.24.1 (bottled), HEAD
Isolated development environments using Docker
https://docs.docker.com/compose/
/usr/local/Cellar/docker-compose/1.24.0 (1,635 files, 17.3MB) *
  Poured from bottle on 2019-06-22 at 23:31:43
From: https://github.com/Homebrew/homebrew-core/blob/master/Formula/docker-compose.rb
=> Dependencies
Required: libyaml ✔, python ✔
=> Options
--HEAD
    Install HEAD version
=> Caveats
Bash completion has been installed to:
  /usr/local/etc/bash_completion.d

zsh completions have been installed to:
  /usr/local/share/zsh/site-functions
=> Analytics
install: 11,264 (30 days), 31,818 (90 days), 125,097 (365 days)
install_on_request: 11,024 (30 days), 31,116 (90 days), 120,355 (365 days)
build_error: 0 (30 days)
MBP-Pawel:~ pawelurbanowicz$

```

## e. Install python 3

```
brew install python
```

```

MBP-Pawel:~ pawelurbanowicz$ python3 --version
Python 3.7.4
MBP-Pawel:~ pawelurbanowicz$

```

## f. Clone code from repository

```
git clone https://github.com/mledl/BDMA\_HW
```

## g. Open terminal and go to docker\_spark\_hadoop directory and run:

```
docker-compose up
```

```

MBP-Pawel:docker_spark_hadoop pawelurbanowicz$ docker-compose up
Creating network "docker_spark_hadoop_default" with the default driver
Creating namenode ... done
Creating spark-master ... done
Creating spark-worker-2 ... done
Creating spark-worker-1 ... done
Creating docker_spark_hadoop_datanode_1 ... done
Attaching to spark-master, namenode, spark-worker-1, spark-worker-2, docker_spark_hadoop_datanode_1
namenode      | Configuring core
spark-master   | Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
namenode      | Setting hadoop proxyuser: hys host: *

```

## h. Go to HW2 directory to build image for python script

```
docker build --rm -t app .
```

```
MBP-Pawel:HW1 pawelurbanowicz$ docker build --rm -t hpc-app .
Sending build context to Docker daemon 96.87MB
Step 1/11 : FROM bde2020/spark-submit:2.4.4-hadoop2.7
----> dac823dd609e
Step 2/11 : COPY /app /app
----> 16b126a91da3
Step 3/11 : COPY /preprocessed /preprocessed
----> a8de2603ec68
Step 4/11 : COPY docker-spark/template.sh /
----> ed53462056b7
Step 5/11 : RUN apk add --update alpine-sdk
----> Running in a83f5ae0b58e
```

It can take some time as some libraries must be built from sources

i. Add data to hadoop

```
docker cp data namenode:data
```

```
docker exec -it namenode bash
```

```
hadoop fs -put /data /data
```

j. Run previously build image

```
docker run -it --name app -e ENABLE_INIT_DAEMON=false --link
```

```
spark-master:spark-master --net docker_spark_hadoop_default -d app
```

```
MBP-Pawel:HW1 pawelurbanowicz$ docker run -it --name hpc-app -e ENABLE_INIT_DAEMON=false --link spark-master:spark-master --net
3f349f1d0aae29ba7228402abd62f8e8fe1d2dfb6623e173d588cdf4e3d1aeb
MBP-Pawel:HW1 pawelurbanowicz$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS
3f349f1d0aae        hpc-app            "/bin/bash /template..."   7 seconds ago       Up 5 seconds
e217b0571e1e        bde2020/spark-worker:2.4.4-hadoop2.7  "/bin/bash /worker.sh"      About an hour ago   Up About an hour   8081
95e8265d6db7        bde2020/spark-worker:2.4.4-hadoop2.7  "/bin/bash /worker.sh"      About an hour ago   Up About an hour   8081
2b841a2810b3        bde2020/spark-master:2.4.4-hadoop2.7  "/bin/bash /master.sh"      About an hour ago   Up About an hour   6066
MBP-Pawel:HW1 pawelurbanowicz$
```

Result of setup:

<http://localhost:8089/>



Spark Master at spark://248fe853406e:7077

URL: spark://248fe853406e:7077

Alive Workers: 2

Cores in use: 8 Total, 8 Used

Memory in use: 2.0 GB Total, 2.0 GB Used

Applications: 1 Running, 2 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (2)

Worker Id	Address	State	Cores	Memory
worker-20191015132226-172.19.0.4-38783	172.19.0.4:38783	ALIVE	4 (4 Used)	1024.0 MB (1024.0 MB Used)
worker-20191015132226-172.19.0.6-45749	172.19.0.6:45749	ALIVE	4 (4 Used)	1024.0 MB (1024.0 MB Used)

<http://localhost:8084/> and <http://localhost:8085/>

## Spark Worker at 172.22.0.6:40427

ID: worker-20191017024419-172.22.0.6-40427

Master URL: spark://54f0b773b528:7077

Cores: 4 (0 Used)

Memory: 2.9 GB (0.0 B Used)

[Back to Master](#)

### Running Executors (0)

ExecutorID	Cores	State	Memory	Job Details	Logs
------------	-------	-------	--------	-------------	------

<http://localhost:9870/>

Hadoop	Overview	Datanodes	Datanode Volume Failures	Snapshot	Startup Progress	Utilities
--------	----------	-----------	--------------------------	----------	------------------	-----------

## Overview 'namenode:8020' (active)

Started:	Tue Oct 15 14:35:02 +0800 2019
Version:	2.8.0, r91f2b7a13d1e97be65db92ddabc627cc29ac0009
Compiled:	Fri Mar 17 12:12:00 +0800 2017 by jdu from branch-2.8.0
Cluster ID:	CID-2f2f1c2f-7b71-4d14-bc34-2d0afe847ef6
Block Pool ID:	BP-1794751137-172.20.0.2-1571100513076

## Summary

### docker ps

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
5f6abc7d5deb	purbanow/spark-worker:latest	"/bin/bash /worker.sh"	5 minutes ago	Up 5 minutes	0.0.0.0:8084->8081/tcp
2c695bdc2856	purbanow/spark-worker:latest	"/bin/bash /worker.sh"	5 minutes ago	Up 5 minutes	0.0.0.0:8085->8081/tcp
3e94d8d3cb4e	bde2020/hadoop-datanode:2.0.0-hadoop3.1.2-java8	"/entrypoint.sh /run..."	5 minutes ago	Up 5 minutes (healthy)	9864/tcp
54f0b773b528	purbanow/spark-master:latest	"/bin/bash /master.sh"	5 minutes ago	Up 5 minutes	6066/tcp, 0.0.0.0:7077->7077/tcp, 0.0.0.0:8089->8089/tcp
9872caf89a8d	bde2020/hadoop-namenode:2.0.0-hadoop3.1.2-java8	"/entrypoint.sh /run..."	5 minutes ago	Up 5 minutes (healthy)	0.0.0.0:9870->9870/tcp

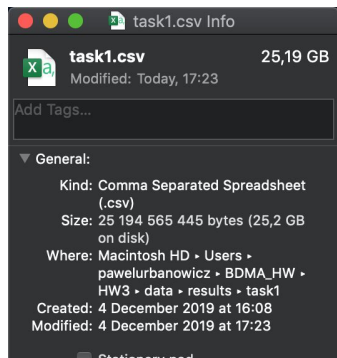
## Data Preprocessing

- Replacing newlines by one whitespace
- Converting all text to lowercase
- removing <BODY> tag
- replacing '?' and '\$' by one whitespace
- replacing multiline whitespaces with one

## Output Format

Task 1: set representation: A  $M \times N$  matrix: with rows as shingles and columns as documents ( $N=21,578$ )

25.12 GB file



Task 2: minhash signatures: The  $H \times N$  signature matrix: with  $H$  as the number of hash functions, and  $N=21,578$

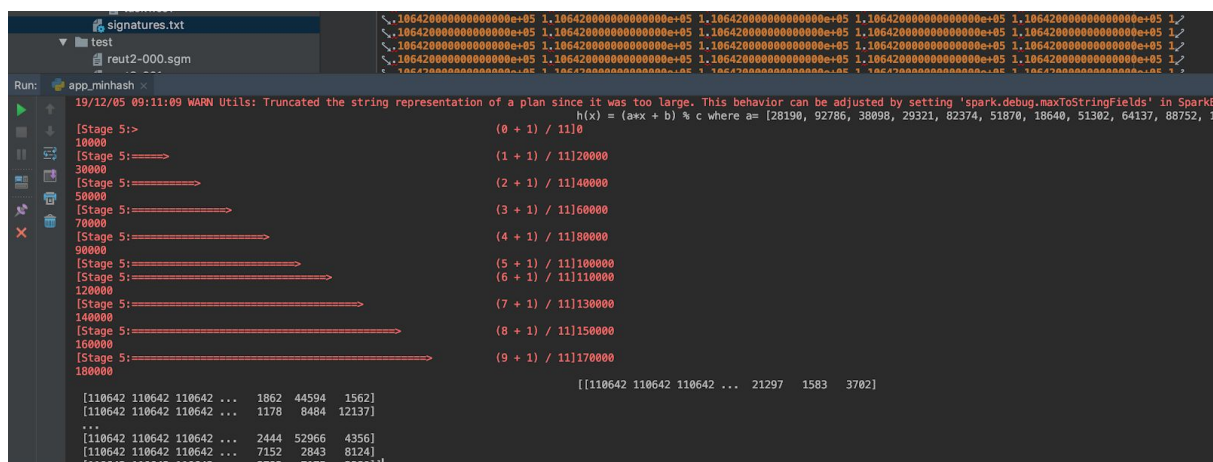
### Spark Jobs (?)

User: pawelurbanowicz  
Total Uptime: 12 min  
Scheduling Mode: FIFO  
Completed Jobs: 5

Event Timeline

Completed Jobs (5)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
4	foreach at /Users/pawelurbanowicz/BDMA_HW/HW3/app/app_minhash.py:97 foreach at /Users/pawelurbanowicz/BDMA_HW/HW3/app/app_minhash.py:97	2019/12/05 09:54:00	4.8 min	1/1	11/11
3	zipWithIndex at /Users/pawelurbanowicz/BDMA_HW/HW3/app/app_minhash.py:96 zipWithIndex at /Users/pawelurbanowicz/BDMA_HW/HW3/app/app_minhash.py:96	2019/12/05 09:50:45	3.2 min	1/1	11/11
2	count at NativeMethodAccessorImpl.java:0 count at NativeMethodAccessorImpl.java:0	2019/12/05 09:50:41	3 s	2/2	12/12
1	csv at NativeMethodAccessorImpl.java:0 csv at NativeMethodAccessorImpl.java:0	2019/12/05 09:49:13	1.4 min	1/1	11/11
0	csv at NativeMethodAccessorImpl.java:0 csv at NativeMethodAccessorImpl.java:0	2019/12/05 09:49:12	0.5 s	1/1	1/1



Task 3: candidate pairs

[illegible]

## Spark Jobs (?)

**User:** pawelurbanowicz

**Total Uptime:** 6.5 min

**Scheduling Mode: FIFO**

Completed Jobs: 5

► [Event Timeline](#)

Completed

▼ Completed Jobs (5)

Job id ↓	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
4	collect at /Users/pawelurbanowicz/BDMA_HW/HW3/app/app_ish.py:61 collect at /Users/pawelurbanowicz/BDMA_HW/HW3/app/app_ish.py:61	2019/12/05 10:02:22	24 s	3/3	201/201
3	count at NativeMethodAccessorImpl.java:0 count at NativeMethodAccessorImpl.java:0	2019/12/05 10:02:22	0.2 s	2/2	2/2
2	showString at NativeMethodAccessorImpl.java:0 showString at NativeMethodAccessorImpl.java:0	2019/12/05 10:02:17	5 s	1/1	1/1
1	csv at NativeMethodAccessorImpl.java:0 csv at NativeMethodAccessorImpl.java:0	2019/12/05 10:02:11	0.5 s	1/1	1/1
0	csv at NativeMethodAccessorImpl.java:0 csv at NativeMethodAccessorImpl.java:0	2019/12/05 10:02:10	0.3 s	1/1	1/1